Executive Summary — Finesse × Citadel Round 2

Strategy. A systematic, long-only, 10-stock portfolio drawn monthly from the Nifty 100 / Midcap 100 / Smallcap 100 universe. We hold the ten most liquid names that best combine **12-month price momentum (55%)** with **low recent volatility (45%)**, weight winners more heavily subject to a 25% cap, and rebalance monthly with a turnover buffer. No leverage, no shorts, 0.1% charged on every trade. It is a rules engine you can read in an afternoon — no black box, no fundamentals we cannot source cleanly, no discretionary overrides.

Headline result (₹1 crore start).

	Backtest 2021–2025	Held-out 2026 H1
Total Net PnL	**₹8.99 crore**	**₹11.93 lakh**
CAGR	60.0%	26.5%
Sharpe (rf = 0%)	2.42	0.93
Max drawdown	−29.4%	−17.1%
vs Nifty 100	+89.4% index	−6.4% index

The out-of-sample half-year is the one that matters: the portfolio made **+11.9% while both the Nifty 100 and Nifty 500 fell**, on a model that was frozen before that data existed.

What makes this credible, not just large. A five-year Indian mid/small-cap bull run flatters almost any long book, so we refuse to let the 899% headline stand on its own. We benchmark against a skill-free equal-weight hold of our *own* universe, which strips out the size, equal-weight and survivorship tailwinds that every team inherits. On that bias-matched basis our **selection alpha is +25.5 pp/yr in-sample and +22.6 pp/yr out-of-sample** — nearly identical across a genuinely held-out window. That stability, not the raw return, is the evidence.

We did not tune to the test. Every parameter is either textbook/mandated or was chosen on a **train (2021–2023) / validate (2024–2025) split that never touched 2026**. Where the in-sample data begged us to push momentum weight to 65–80%, the validation set said 50–55% generalises better — so we took 55% and left the higher in-sample number on the table. The 2026 result was computed only after the model was frozen. This is the "defensible process over hindsight" the brief asks for, made literal.

Three stress tests we ran on ourselves.

1. *Zero-skill controls* — random and alphabetical 10-stock portfolios score ≈0 selection alpha in-sample, confirming our measurement isn't crediting mere concentration.
2. *Point-in-time universe* — on the August-2023 constituent lists (78% of the Smallcap 100 has since changed), with no survivorship advantage, the strategy still earns **+13.5 pp/yr selection alpha in a falling 2026 market**.

3. *14 alternative signals* — reversal, trend, low-beta, illiquidity, residual momentum, sector-neutral and ensemble variants. None beats the shipped rule on the bar that matters (positive every year *and* out-of-sample), and reversal's failure confirms our momentum signal has the right sign.

Known limitations, stated plainly. The universe snapshot carries survivorship bias (quantified, ~+20 pp/yr, common to all teams); the out-of-sample window is short (trust the stability of the edge, not the single level); 0.1% models explicit cost only; the book is price-only by design. All are discussed in the report and none changes the conclusion.

Bottom line. ₹1 crore becomes ₹9.99 crore over the backtest and grows a further ₹11.93 lakh out-of-sample while the market falls — but the number we would defend to a jury is the +22.6 pp/yr of out-of-sample selection alpha over our own universe, produced by a rule that was frozen before it ever saw that data.